

COMP90054 – Week 9 tutorial

Last updated: 1 May 2023

Multi-armed Bandits (MABs)

Definition

- A stochastic multi-armed bandit is a set of random variables $X_{i,t}$ where each variable represents a reward from playing arm i at time t
 - $i \in \{1, \dots, N\}$ is the arm of the bandit
 - $k \in \{1, \dots, K\}$ represents the index of play
- Successive plays are assumed to be independent from each other
- Each arm provides a reward drawn from an unknown probability distribution
- MAB in RL Context: At state s , the agent has a range of “arms” (a set of actions $A(s)$) that it can choose from, but the reward it receives from playing an arm is uncertain

MAB Strategies

- Aim: To learn $\pi(k)$, i.e. learn which arm to play at the k -th play to maximise the sum of the rewards collected over all rounds
- Problem: want to only play good actions, but don't know which ones are good
- Aim: To learn $\pi(k)$, i.e. learn which arm to play at the k -th play to maximise the sum of the rewards collected over all rounds
- Problem: want to only play good actions, but don't know which ones are good
- Intuition: find balance between
 - exploitation: play the arms deemed as the best so far
 - exploration: play other arms to potentially identify better arm

Ideas of some strategies

- Epsilon-greedy: Explore with probability of ϵ , exploit with probability of $1 - \epsilon$
- Epsilon-decreasing: Explore less often as we play more rounds (ϵ decays)
 - Softmax: Probability of choosing a depends on the empirical expected reward of a , can be controlled by the temperature parameter τ
- Upper Confidence Bounds (UCB1): An asymptotically optimal approach that considers both the empirical mean reward of each arm and the uncertainty about the true mean reward

Temporal difference reinforcement learning

Model-free reinforcement learning

- Idea: To learn a *policy* from the simulated *episodes* of the MDP problem
 - e.g., learn $Q(s, a)$ for all state-action pairs, then extract policy
- During each *episode*, starting from initial state s_0 , a simulator generates a sequence of actions, rewards, and subsequent states.
 - At each state s , an action a is chosen, and the reward r and the new state s' are observed
 - Model-free RL techniques *reinforce* the Q estimates of applying a in s with the new observation
- Terminate the process if:
 - we run out of training time
 - we think our policy has converged to the optimal policy (for each new episode we see no improvement)
 - our policy is 'good enough' (for each new episode we see minimal improvement)

Model-based RL vs Model-free RL

- Similarity: both work with MDPs
- Difference: whether they have/use knowledge on $P_a(s'|s)$ and $r(s, a, s')$
 - Model-based RL uses this knowledge, e.g. value iteration
 - Model-free RL does not, and they rely on a simulator and learn through experience
- Given s and a , the simulator provides observations of s' according to $P_a(s'|s)$ and $r(s, a, s')$
- Model-free RL can use (but doesn't have to) use MAB strategies to select actions

Temporal difference (TD) methods

- TD methods use bootstrapping to alleviate the high variance issue in learning $Q(s, a)$ via Monte-Carlo RL
- Bootstrapping: update value function with agent's estimates rather than waiting for actual discounted rewards to be realised
- TD methods update $Q(s, a)$ by

$$Q(s, a) \leftarrow \underbrace{Q(s, a)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \left[\underbrace{r + \gamma \cdot V(s')}_{\text{TD target}} - \underbrace{Q(s, a)}_{\text{do not count extra } Q(s, a)} \right]$$

- - New value of $Q(s, a)$ is a weighted average between TD target and the old value weighted by learning rate α
- Q-learning and SARSA are two examples of TD methods

Q-learning

- Idea: Estimate $V(s')$ with $\max_{a' \in A(s')} Q(s', a')$
 - Execute a , Observe reward r and new state s'
 - $Q(s, a) \leftarrow Q(s, a) + \alpha \cdot \left[r + \gamma \cdot \max_{a' \in A(s')} Q(s', a') - Q(s, a) \right]$

SARSA

- State-Action-Reward-State-Action
 - $s - a - r - s' - a'$
- Idea: Estimate $V(s')$ with the actual $Q(s', a')$ obtained in simulation
 - Execute a , observe reward r and new state s'
 - Select a' to apply in s' (thereby determining both state and action for next step)
 - $Q(s, a) \leftarrow Q(s, a) + \alpha \cdot [r + \gamma \cdot Q(s', a') - Q(s, a)]$

Comparing Q-learning and SARSA

- Q-learning is off-policy
 - it uses the greedy reward estimate for next state in its update rather than following the policy
 - it will converge to the optimal policy (assuming that all state-action pairs are sampled infinitely often), irrespective of the policy followed
- SARSA is on-policy
 - it chooses its action using the same policy used to choose the previous action, and then uses this difference to update its Q-function
 - its learning is dependent on the policy used